

finalprob

May 15, 2023

1 Smit Kumbhani 114964474

2 Mitesh Jalan 115222506

3 Pritish Thombare 115382224

```
[133]: from google.colab import drive
drive.mount('/content/drive')
import numpy as np
import pandas as pd
import math
import matplotlib.pyplot as plt
data = pd.read_csv('/content/drive/My Drive/ProbStatsProject/
↳OriginalDallas97-22.csv', sep=',')
print(data)
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

	DATE	CPI_all_items	Rent_of_Primary_residence \
0	1997-01-01	159.400	140.400
1	1997-02-01	159.700	140.500
2	1997-03-01	159.800	140.500
3	1997-04-01	159.900	141.000
4	1997-05-01	159.900	140.800
..
307	2022-08-01	295.320	326.759
308	2022-09-01	296.539	330.192
309	2022-10-01	297.987	333.473
310	2022-11-01	298.598	338.266
311	2022-12-01	298.990	341.333

	Monthly_Housing_Cost	CPI_Energy	US_Dollar_Purchasing_power
0	.	126.000	62.8
1	137.9	129.200	62.6
2	.	118.200	62.5
3	136.7	107.200	62.4
4	.	108.200	62.5

..
307	.	331.995	33.8
308	276.995	336.048	33.7
309	.	338.534	33.6
310	280.344	310.256	33.6
311	.	311.458	33.7

[312 rows x 6 columns]

3.1

```
[134]: #Tukey
cols = list(data.columns)
#cols withou DATE
cpi_cols = cols[1:]
alpha = 1.5
def getLimits(values):
    values.sort()
    length = len(values)
    q1 = values[math.ceil((1/4)*length)-1]
    q3 = values[math.ceil((3/4)*length)-1]
    IQR = q3-q1
    right = q3 + alpha*IQR
    left = q1 - alpha*IQR
    return left, right

def removeOutliers(data, col):
    clist = data[col].tolist()
    clist = [float(c) for c in clist if c != '.']
    left, right = getLimits(clist)
    print('Limits : ', round(left,2), round(right,2))
    df = data.apply(lambda x : True if (x[col] != '.' and (float(x[col]) > right_
    or float(x[col]) < left)) else False, axis = 1)
    outliers = len(df[df == True].index)
    print('Outliers in the column: ', col, outliers)
    data[col] = data.apply(lambda x: '.' if (x[col] != '.' and (float(x[col]) >
    right or float(x[col]) < left)) else x[col], axis=1)

for col in cpi_cols:
    removeOutliers(data, col)
#count the number of outliers in each column
```

Limits : 99.1 324.7

Outliers in the column: CPI_all_items 0

Limits : 73.28 331.74

Outliers in the column: Rent_of_Primary_residence 3

Limits : 102.46 264.97

Outliers in the column: Monthly_Housing_Cost 4
Limits : 76.45 306.18
Outliers in the column: CPI_Energy 6
Limits : 22.3 73.5
Outliers in the column: US_Dollar_Purchasing_power 0

3.2

[135]: *#Interpolation*

```
edgevals = {'CPI_all_items':[159.400, 298.900], 'Rent_of_Primary_residence':
    ↳[164.400, 385.649], 'Monthly_Housing_Cost':[155.100, 310.725], 'CPI_Energy':
    ↳[115.200, 287.176], 'US_Dollar_Purchasing_power':[62.8, 33.7]}
def interpolation(data, col, edge):
    ls = data[col].tolist()
    ls.append(edge[col][1])
    ls.insert(0,edge[col][0])
    #print(ls)

    for i in range(1,len(ls)-1):
        if ls[i] == '.':
            up = i-1
            down = -1
            #getting the next know value
            for j in range(i, len(ls)):
                if ls[j] != '.':
                    down = j
                    break
            slope = (float(ls[down]) - float(ls[up]))/(down-up)
            missing = float(ls[up]) + slope*(i-up)
            ls[i] = round(missing,3)

    ls = ls[1:-1]
    ls = [float(v) for v in ls]
    data[col] = ls

for col in cpi_cols:
    interpolation(data, col, edgevals)

print(data)
```

	DATE	CPI_all_items	Rent_of_Primary_residence \
0	1997-01-01	159.400	140.400
1	1997-02-01	159.700	140.500
2	1997-03-01	159.800	140.500
3	1997-04-01	159.900	141.000
4	1997-05-01	159.900	140.800
..

307	2022-08-01	295.320	326.759
308	2022-09-01	296.539	330.192
309	2022-10-01	297.987	344.056
310	2022-11-01	298.598	357.920
311	2022-12-01	298.990	371.784

	Monthly_Housing_Cost	CPI_Energy	US_Dollar_Purchasing_power
0	146.500	126.000	62.8
1	137.900	129.200	62.6
2	137.300	118.200	62.5
3	136.700	107.200	62.4
4	137.850	108.200	62.5
..
307	286.217	297.026	33.8
308	291.119	295.056	33.7
309	296.021	293.086	33.6
310	300.922	291.116	33.6
311	305.824	289.146	33.7

[312 rows x 6 columns]

3.4 KS Test and Permutation Test

[136]: *#Normalizing the data*

```
data['CPI_all_items_per_change'] = data['CPI_all_items'].pct_change()
data['Rent_of_Primary_residence_per_change'] =
    ↪data['Rent_of_Primary_residence'].pct_change()
data['Monthly_Housing_Cost_per_change'] = data['Monthly_Housing_Cost'].
    ↪pct_change()
data['CPI_Energy_per_change'] = data['CPI_Energy'].pct_change()
data['US_Dollar_Purchasing_power_per_change'] =
    ↪data['US_Dollar_Purchasing_power'].pct_change()

#Removing negative values and interpolating them for some distributions
#Replacing negative values with a '.'
data['Monthly_Housing_Cost_per_change_pos'] = data.apply(lambda x: '.' if
    ↪(x['Monthly_Housing_Cost_per_change'] < 0 or np.
    ↪isnan(x['Monthly_Housing_Cost_per_change'])) else
    ↪x['Monthly_Housing_Cost_per_change'], axis=1)
data['CPI_all_items_per_change_pos'] = data.apply(lambda x: '.' if
    ↪(x['CPI_all_items_per_change'] < 0 or np.
    ↪isnan(x['CPI_all_items_per_change'])) else x['CPI_all_items_per_change'],
    ↪axis=1)
data['Rent_of_Primary_residence_per_change_pos'] = data.apply(lambda x: '.' if
    ↪(x['Rent_of_Primary_residence_per_change'] < 0 or np.
    ↪isnan(x['Rent_of_Primary_residence_per_change'])) else
    ↪x['Rent_of_Primary_residence_per_change'], axis=1)
```

```

data['CPI_Energy_per_change_pos'] = data.apply(lambda x: '.' if
↳(x['CPI_Energy_per_change'] < 0 or np.isnan(x['CPI_Energy_per_change']))
↳else x['CPI_Energy_per_change'], axis=1)
data['US_Dollar_Purchasing_power_per_change_pos'] = data.apply(lambda x: '.' if
↳(x['US_Dollar_Purchasing_power_per_change'] < 0 or np.
↳isnan(x['US_Dollar_Purchasing_power_per_change'])) else
↳x['US_Dollar_Purchasing_power_per_change'], axis=1)

#Interpolating the per_change data to get positive values.
#Placing 0.0 before and after the list of values for interpolation.
edgevals_pos = {'CPI_all_items_per_change_pos':[0.0, 0.0],
↳'Rent_of_Primary_residence_per_change_pos':[0.0, 0.0],
↳'Monthly_Housing_Cost_per_change_pos':[0.0, 0.0],
↳'CPI_Energy_per_change_pos':[0.0, 0.0],
↳'US_Dollar_Purchasing_power_per_change_pos':[0.0, 0.0]}
for col in cpi_cols:
    col = col + '_per_change_pos'
    interpolation(data, col, edgevals_pos)

data.fillna(0, inplace=True)

```

1 Sample KS Test

```

[137]: from scipy.stats import poisson
from scipy.stats import geom
from scipy.stats import binom

data_18_20 = pd.DataFrame()
for index, row in data.iterrows():
    if '2018' in row['DATE'] or '2019' in row['DATE'] or '2020' in row['DATE']:
        data_18_20 = data_18_20.append(row, ignore_index=True)

x_bar = data_18_20['Monthly_Housing_Cost'].mean()
variance = data_18_20['Monthly_Housing_Cost'].var()
rent = sorted(data_18_20['Rent_of_Primary_residence']).tolist()

#poisson cdf
def poisson_cdf(x):
    lambda_mme = x_bar
    poisson_dist = poisson(lambda_mme)
    return poisson_dist.cdf(x)

#binomial pdf

```

```

def binomial_cdf(x):
    p_mme = 1 - (variance/x_bar)
    n_mme = x_bar/p_mme
    binomial_dist = binom(n_mme, p_mme)
    return binomial_dist.cdf(x)

#geometric pdf
def geometric_cdf(x):
    pmme = 1/x_bar
    geometric_dist = geom(pmme)
    return geometric_dist.cdf(x)

rent_ecdf = {}
for r in rent:
    count = sum(1 for x in rent if x <= r)
    rent_ecdf[r] = count/len(rent)

def one_sampleks(X, F_X):
    maxvalue = 0
    for i in range(0, len(X)):
        FDminus, FDplus = 0, 1
        if i != 0:
            FDminus = rent_ecdf[X[i-1]]
        if i != len(X)-1:
            FDplus = rent_ecdf[X[i]]
        max_temp = max(abs(F_X[i] - FDplus), abs(F_X[i] - FDminus))
        if max_temp > maxvalue:
            maxvalue = max_temp
    return maxvalue

poisson = [poisson_cdf(x) for x in rent]
geometric = [geometric_cdf(x) for x in rent]
binomial = [binomial_cdf(x) for x in rent]
#approximating binomial cdf
binomial = [1 if math.isnan(x) else x for x in binomial]

print('KS Test for poisson is : ', one_sampleks(rent, poisson))
print('KS Test for geometric is : ', one_sampleks(rent, geometric))
print('KS Test for binomial is : ', one_sampleks(rent, binomial))

```

<ipython-input-137-db0b490d122c>:9: FutureWarning: The frame.append method is deprecated and will be removed from pandas in a future version. Use pandas.concat instead.

```
data_18_20 = data_18_20.append(row, ignore_index=True)
```

```
KS Test for poisson is : 0.9709350135645295
```

```
KS Test for geometric is : 0.6758968422271779
```

```
KS Test for binomial is : 1
```

The KS statistic values are greater than the threshold of 0.05 which means that we reject H_0 . H_0 is that the sample is from the given distribution. We can conclude that the sample is not from Poisson, Geometric and Binomial distribution.

```
[138]: #2 sample KS test
#Between Rent and House data
house = sorted(data_18_20['Monthly_Housing_Cost_per_change'].tolist())
rent = sorted(data_18_20['Rent_of_Primary_residence_per_change'].tolist())

def eCDF(ls):
    ls.sort()
    p = 1/(len(ls))
    x = []
    f_x = []
    tot = p
    for i in range(0,len(ls)):
        x = x + [ls[i]]
        f_x = f_x + [tot]
        tot += p
    return x, f_x

def two_sample_ks_test(x, y, X, f_X, Y, f_Y):
    i=0
    j=0
    maxi=0
    while i<len(x) and j<len(y):
        if x[i] < y[j]:
            if j==0:
                F_x_cap=0
            else:
                F_x_cap = f_Y[Y.index(y[j-1])]
            if i==0:
                F_y_left=0
            else:
                F_y_left = f_X[X.index(X[i-1])]
            if i==len(X)-1:
                F_y_right=1
            else:
                F_y_right=f_X[X.index(x[i])]

            maxi_temp=max(abs(F_x_cap-F_y_left),abs(F_x_cap-F_y_right))
            if maxi_temp > maxi:
```

```

        maxi = maxi_temp
        index = i
        i+=1
    else:
        j+=1
    return maxi

rent.sort()
house.sort()
X, f_X = eCDF(rent)
Y, f_Y = eCDF(house)


KS_p_value = two_sample_ks_test(rent, house, X, f_X, Y, f_Y)

print('KS statistic : ', KS_p_value)
plt.figure('K-S test' , figsize=(10,7))
plt.xlabel("K-S 2 Sample Test")
plt.ylabel('eCDF')

plt.step(X, f_X, where='post', label="Rent")
plt.step(Y, f_Y, where='post', label="Housing")
plt.legend()
plt.show()

```

KS statistic : 0.3333333333333334



The KS statistic is 0.333 which is greater than 0.05, thus we can reject the null hypothesis that the two samples are from the same distribution.

Permutation Test

```
[139]: def calculate_mean(arr):
    lst = arr.to_list()
    return sum(lst) / len(lst)

rent_mean = calculate_mean(data_18_20['Rent_of_Primary_residence_per_change'])
housing_mean = calculate_mean(data_18_20['Monthly_Housing_Cost_per_change'])
Tobs = abs(rent_mean - housing_mean)

def permutations(n, d1, d2, Tobs):
    d1arr = np.array(d1.to_list())
    d2arr = np.array(d2.to_list())
    darr = np.append(d1arr, d2arr, axis=0)
    ↪ #appending the two arrays
    ld1 = len(d1arr)
    ld2 = len(d2arr)
    extremes = 0
    ↪ #calculates the cases where  $T_i > T_{obs}$ 
```

```

for i in range(0, n):
    ↪#runs for the number of random permutations
    np.random.shuffle(darr)
    new_d1 = darr[0:ld1]
    new_d2 = darr[ld1:]
    Ti = abs(np.mean(new_d1) - np.mean(new_d2))
    if Ti > Tobs:
        extremes += 1
    return extremes / n

print('P value is : ', permutations(1000,
    ↪data_18_20['Rent_of_Primary_residence_per_change'],
    ↪data_18_20['Monthly_Housing_Cost_per_change'], Tobs))

```

P value is : 0.363

The P value comes out to be 0.366 which is greater than 0.05 thus we accept the null hypothesis that the samples are from the same distribution.

Normality Testing

QQ Plot

```

[140]: import seaborn as sns
import statsmodels.api as sm
from scipy.stats import truncnorm

def get_truncated_normal(mean=0, sd=1, low=0, upp=10):
    return truncnorm(
        (low - mean) / sd, (upp - mean) / sd, loc=mean, scale=sd)

data_16_21 = pd.DataFrame()
for index, row in data.iterrows():
    if '2016' in row['DATE'] or '2017' in row['DATE'] or '2018' in row['DATE']
    ↪or '2019' in row['DATE'] or '2020' in row['DATE'] or '2021' in row['DATE']:
        data_16_21 = data_16_21.append(row, ignore_index=True)
    #print(data_18_20)

# Generating samples from normal distribution
sample_size = len(data_16_21['CPI_Energy_per_change'])
energy = data_16_21['CPI_Energy_per_change'].tolist()
all_items = data_16_21['CPI_all_items_per_change'].tolist()

# Sorting the samples
energy_sorted = np.sort(energy)
all_items_sorted = np.sort(all_items)

#Getting truncated normal values within the range of the quantile values of the
    ↪data.

```

```

X = get_truncated_normal(mean=0, sd=1, low=-0.005, upp=0.01)
scnorm = X.rvs(sample_size)
scnsorted = np.sort(scnorm)

# Calculate quantiles
p = np.arange(1, sample_size+1) / sample_size
q2 = np.quantile(energy_sorted, p)
q3 = np.quantile(all_items_sorted, p)
q4 = np.quantile(scnsorted, p)

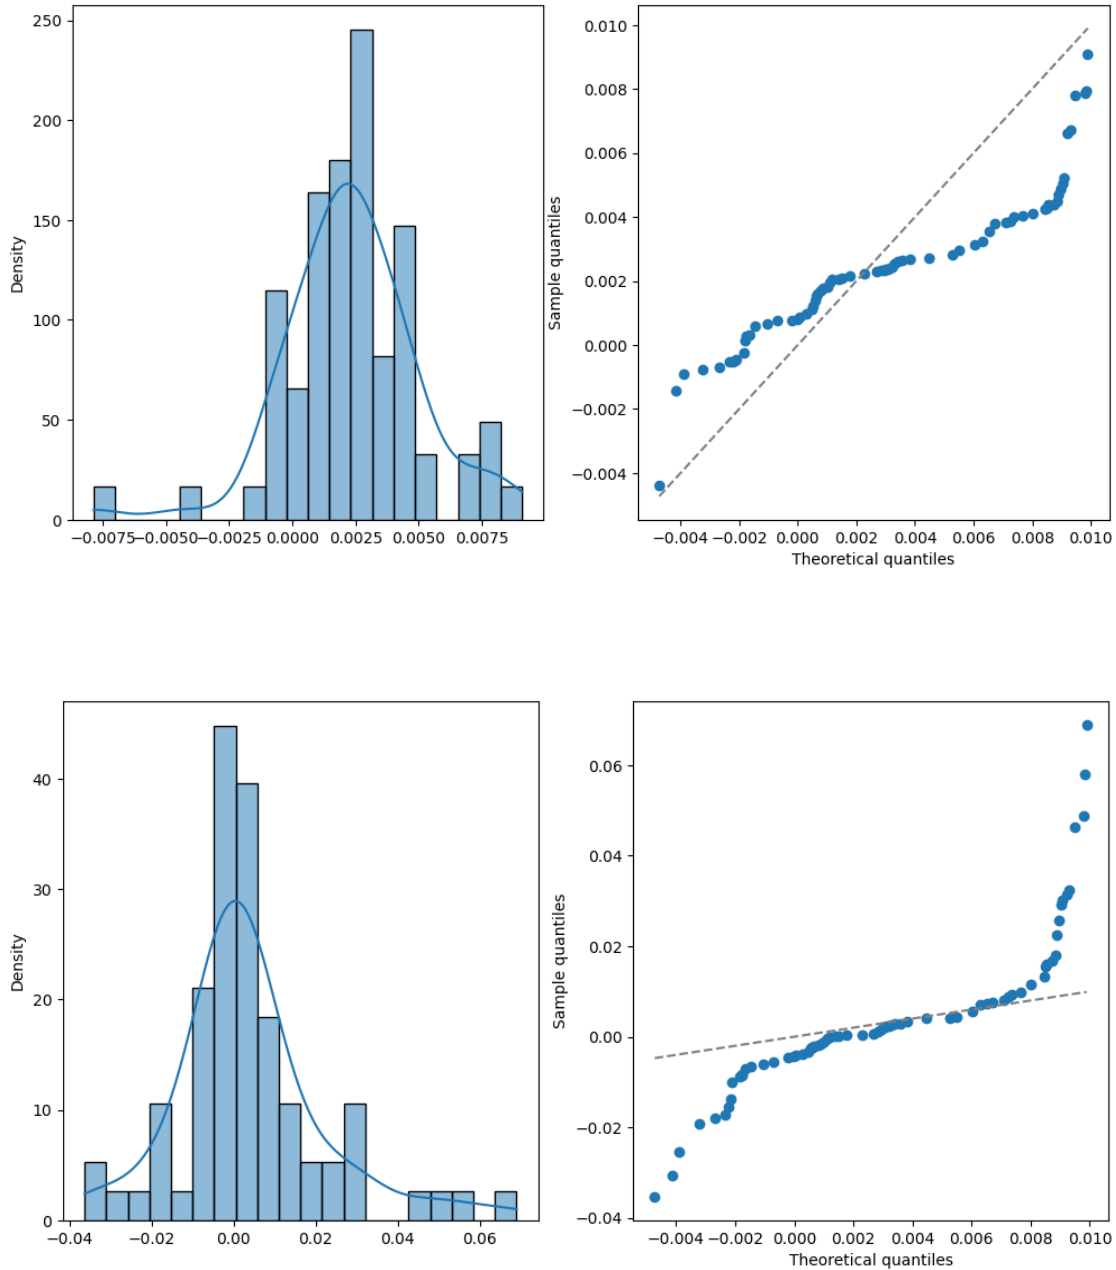
# Plotting Q-Q plot for column 'All items'
fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(12, 6))
sns.histplot(all_items_sorted, ax=ax1, bins=20, kde=True, stat='density')
plt.scatter(q4, q3)
ax2.plot([q4.min(), q4.max()], [q4.min(), q4.max()], '--', color='gray')
ax2.set_xlabel('Theoretical quantiles')
ax2.set_ylabel('Sample quantiles')
plt.show()

# Plotting Q-Q plot for column 'Energy'
fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(12, 6))
sns.histplot(energy_sorted, ax=ax1, bins=20, kde=True, stat='density')
plt.scatter(q4, q2)
ax2.plot([q4.min(), q4.max()], [q4.min(), q4.max()], '--', color='gray')
ax2.set_xlabel('Theoretical quantiles')
ax2.set_ylabel('Sample quantiles')
plt.show()

```

<ipython-input-140-4c9d4a1e9125>:12: FutureWarning: The frame.append method is deprecated and will be removed from pandas in a future version. Use pandas.concat instead.

```
data_16_21 = data_16_21.append(row, ignore_index=True)
```



We have used truncated Normal distribution as Theoretical quantiles, this is to clearly show the distribution through the scatter plot.

We observe that the data is not actually normally distributed as the scatter plot does not align with the normal line.

Shapiro Test

```
[141]: import scipy.stats as stats

df = pd.read_csv("/content/drive/My Drive/ProbStatsProject/PerChangeDallas97-22.
↳csv")
df = df.fillna(0)
all_items_col = df['CPI_all_items_per_change']
energy_col = df['CPI_Energy_per_change']

res1 = stats.shapiro(all_items_col)
print("Column: CPI_all_items_per_change' --> W statistic: "+str(res1[0])+", P-Value: "+str(res1[1]))
res2 = stats.shapiro(energy_col)
print("Column: CPI_Energy_per_change --> W statistic: "+str(res2[0])+", P-Value: "+str(res2[1]))
```

Column: CPI_all_items_per_change' --> W statistic: 0.9080091714859009, P-Value: 7.239740491132851e-13

Column: CPI_Energy_per_change --> W statistic: 0.8547593355178833, P-Value: 1.7685022234710225e-16

The p-values obtained above are less than 0.05 significance level. Thus we reject the null hypothesis that the distributions are normal. This means that the distributions of the two columns are not normal.

3.3 Performing Wald's test, Z-test and t-test

In this step, we want to check how the mean of monthly stats has changed between 2020 and 2021 (if your dataset is from 1997-2022) and 1994 and 1995 (if your dataset is from 1970-1996). Apply the Wald's test, Z-test, and t-test (assume all are applicable) to check whether the mean of Consumer Price Index for All Urban Consumers: Rent of Primary Residence in 'assigned urban center' and Consumer Price Index for All Urban Consumers: Energy in 'assigned urban center' are different for given years in the urban center. Do this separately for both columns, i.e., you have to compare mean of monthly stats from year 1 with mean of monthly stats from year 2 separately for both columns. Use MLE for Wald's test as the estimator; assume for Wald's estimator purposes that daily data is Poisson distributed.

```
[142]: import pandas as pd
import numpy as np
```

```
[143]: data=pd.read_csv("/content/drive/My Drive/ProbStatsProject/CleanedDallas97-22.
↳csv",sep=",")
```

```
[144]: data.head()
```

```
[144]:
```

	Unnamed: 0	DATE	CPI_all_items	Rent_of_Primary_residence	\
0	0	1997-01-01	159.4	140.4	
1	1	1997-02-01	159.7	140.5	
2	2	1997-03-01	159.8	140.5	
3	3	1997-04-01	159.9	141.0	

4	4	1997-05-01	159.9	140.8
		Monthly_Housing_Cost	CPI_Energy	US_Dollar_Purchasing_power
0		146.50	126.0	62.8
1		137.90	129.2	62.6
2		137.30	118.2	62.5
3		136.70	107.2	62.4
4		137.85	108.2	62.5

```
[145]: start_date = '2020-01-01'
end_date = '2021-12-31'
# Select DataFrame rows between two dates
mask = (data['DATE'] > start_date) & (data['DATE'] <= end_date)
data_3_3=data.loc[mask]
# Rent_of_Primary_residence and CPI_Energy
# (data_3_3.DATE.year==2021).Rent_of_Primary_residence.mean()
df2021=data_3_3.loc[(data_3_3["DATE"]>"2020-12-31")]
# df2021
df2020=data_3_3.loc[(data_3_3["DATE"]<="2020-12-31")]
# df2020
```

Since Poisson distribution is given we will have lambda equal to the mean

```
[146]: #lambda1 - 2020 and lambda2 - 2021

lambda1_energy=df2020.CPI_Energy.mean()
lambda1_rent=df2020.Rent_of_Primary_residence.mean()
X_bar1_energy=lambda1_energy
Y_bar1_rent=lambda1_rent
lambda2_rent=df2021.Rent_of_Primary_residence.mean()
lambda2_energy=df2021.CPI_Energy.mean()
X_bar2_energy=lambda2_energy
Y_bar2_rent=lambda2_rent
```

3.1 Walds test

Given poisson distribution lambda mle= mean

walds test= $\{X_bar-Y_bar\}/\sqrt{\text{var}(X_bar-Y_bar)}$

$\sqrt{\text{var}(X_bar-Y_bar)}$ = $\sqrt{\text{var}(X1)(1/n)+\text{var}(X2)(1/m)}$ = $\sqrt{\text{lambda1}(1/n)+\text{lambda2}(1/m)}$

H0 : lambda1==lambda2

```
[147]: def Walds(X_bar,Y_bar,n,m):
    res=(X_bar-Y_bar)/np.sqrt((X_bar/n)+(Y_bar/m))
    return res
```

```
[148]: walds_energy=Walds(lambda1_energy,lambda2_energy,df2020.CPI_Energy.  
↪shape[0],df2021.CPI_Energy.shape[0])
```

```
[149]: walds_energy
```

```
[149]: -2.6417546664051983
```

$|Walds_energy| > 1.96$ reject H_0 ; here we reject null hypothesis hence $\lambda_1 \neq \lambda_2$

so the mean changes for CPI_Energy

p- value= $2(1-\phi(|w|))$ i.e around 0.00836 which is less than 0.05 value so rejecting null hypothesis

```
[150]: walds_rent=Walds(lambda1_rent,lambda2_rent,df2020.Rent_of_Primary_residence.  
↪shape[0],df2021.Rent_of_Primary_residence.shape[0])
```

```
[151]: walds_rent
```

```
[151]: -1.312849226097177
```

$|Walds_rent| < 1.96$ accept H_0 ; here we accept null hypothesis hence $\lambda_1 = \lambda_2$

So the mean remain same for CPI Rent

p- value= $2(1-\phi(|w|))$ i.e around 0.1902 which is greater than 0.05 value so accept null hypothesis

3.2 Z-Test

variance is known $z = (\bar{x}_1 - \bar{x}_2) / \sqrt{((1)^2/n_1 + (2)^2/n_2)}$

```
[152]: def sigma_square(x):  
    res,s=0,0  
    #summation of xi and x bar:  
    for i in range(len(x)):  
        s+=((x[i]-np.mean(x))**2)  
    return (1/(len(x)-1))*s
```

```
[153]: def Z_test(X_bar,Y_bar,n,m):  
    res=(X_bar-Y_bar)/np.sqrt((X_bar/n)+(Y_bar/m))  
    return res
```

```
[154]: Z_test_energy=Z_test(lambda1_energy,lambda2_energy,df2020.CPI_Energy.  
↪shape[0],df2021.CPI_Energy.shape[0])  
Z_test_energy
```

```
[154]: -2.6417546664051983
```

$|Z_Test| > 1.96$ so null hypothesis is rejected mean of energy changes with year

```
[155]: Z_test_Rent=Z_test(lambda1_rent,lambda2_rent,df2020.Rent_of_Primary_residence.  
↪shape[0],df2021.Rent_of_Primary_residence.shape[0])  
Z_test_Rent
```

```
[155]: -1.312849226097177
```

$|Z\text{-Test}| < 1.96$ so null hypothesis is accepted mean of rent is same with year

3.3 T- Test

<https://libguides.library.kent.edu/spss/independenttttest#:~:text=The%20Independent%20Samples%20t%20Test>

```
[156]: def s_square(x):  
    res,s=0,0  
    #summation of xi and x bar:  
    for i in range(len(x)):  
        s+=((x[i]-np.mean(x))**2)  
    return (1/(len(x)-1))*s
```

```
[157]: def T_Test(X_bar,Y_bar,df1,df2):  
    res=(X_bar-Y_bar)/np.sqrt((s_square(df1.to_numpy())/df1.  
↪shape[0])+(s_square(df2.to_numpy())/df2.shape[0]))  
    return res
```

```
[158]: T_test_energy=T_Test(lambda1_energy,lambda2_energy,df2020.CPI_Energy,df2021.  
↪CPI_Energy)  
T_test_energy
```

```
[158]: -3.0081906840218235
```

if $|T_test_energy| > t(0.025,11) / 2.200985$

we reject Null hypothesis. hence mean for energy is different in 2020 and 2021

```
[159]: T_test_Rent=T_Test(lambda1_rent,lambda2_rent,df2020.  
↪Rent_of_Primary_residence,df2021.Rent_of_Primary_residence)  
T_test_Rent
```

```
[159]: -6.082768916719216
```

if $|T_test_energy| > t(0.025,11) / 2.200985$

we reject Null hypothesis. hence mean for energy is different in 2020 and 2021

```
[160]: import pandas as pd  
import numpy as np  
data1=pd.read_csv("/content/drive/My Drive/ProbStatsProject/  
↪PerChangeDallas97-22.csv")  
data1.columns
```



```
import seaborn as sns
```

```
[161]: import matplotlib.pyplot as plt
```


```
[162]: def s(x,y):  
        return sns.regplot(x,y,ci=None,color='red')
```

```
[163]: def Linear1(X,Y):  
        X_T=X.T  
        B = np.linalg.inv(X_T @ X) @ X_T @ Y  
        B.index=X.columns  
        print(B)  
        pred=X @ B  
        B.index=X.columns  
        print("SSE:",((Y - pred) ** 2).sum())  
        plt.plot(X.iloc[:,1],pred,color='r',label="Predicted")  
        plt.scatter(X.iloc[:,1],Y,color='b',label="Actual")  
        plt.legend()  
        # s(Y,pred)  
        # sns.regplot(x=Y,y=pred,ci=None,color='red')
```

1. US Dollar Purchasing Power per change (CPI All items per change vs US Dollar Purchasing Power per change)

```
[164]: data1["intercept"]=1  
data1=data1.fillna(0)  
Linear1(data1.loc[:  
    ↳,['intercept','US_Dollar_Purchasing_power_per_change']],data1.loc[:  
    ↳,['CPI_all_items_per_change']])
```

	CPI_all_items_per_change
intercept	0.000745
US_Dollar_Purchasing_power_per_change	-0.643533
SSE: CPI_all_items_per_change	0.000907
dtype:	float64




2. CPI Energy per change (CPI All items per change vs [US Dollar Purchasing Power per change, CPI Energy per change])

```
[165]: def Linear2(X,Y):
    X_T=X.T
    B = np.linalg.inv(X_T @ X) @ X_T @ Y
    B.index=X.columns
    # print(B)
    pred=X @ B
    B.index=X.columns
    print("SSE:",((Y - pred) ** 2).sum())
    fig, axes = plt.subplots(2, 1, figsize=(15, 5), sharey=True)
    plt.subplot(2, 1, 1)
    plt.plot(X.iloc[:,1],pred,color='r',label="Predicted")
    plt.scatter(X.iloc[:,1],Y,color='b',label="Actual")
    plt.subplot(2, 1, 2)
    plt.plot(X.iloc[:,2],pred,color='r',label="Predicted")
    plt.scatter(X.iloc[:,2],Y,color='b',label="Actual")
    # sns.regplot(x=Y,y=pred,ci=None,color='red')
    plt.legend()
```

```
[166]: Linear2(data1.loc[:
↳,['intercept','US_Dollar_Purchasing_power_per_change','CPI_Energy_per_change']],data1.
↳loc[:,['CPI_all_items_per_change']])
```

SSE: CPI_all_items_per_change 0.000906
dtype: float64



3. Monthly Housing Cost per change (CPI All items per change vs [US Dollar Purchasing Power per change, CPI Energy per change, Monthly Housing Cost per change])

```
[167]: def Linear3(X,Y):
X_T=X.T
B = np.linalg.inv(X_T @ X) @ X_T @ Y
B.index=X.columns
# print(B)
pred=X @ B
B.index=X.columns

# print(X)
print("SSE:",((Y - pred) ** 2).sum())
sns.regplot(x=Y,y=pred,ci=None,color='red')
plt.subplot(3, 1, 1)
plt.plot(X.iloc[:,1],pred,color='r',label="Predicted")
plt.scatter(X.iloc[:,1],Y,color='b',label="Actual")
plt.subplot(3, 1, 2)
plt.plot(X.iloc[:,2],pred,color='r',label="Predicted")
plt.scatter(X.iloc[:,2],Y,color='b',label="Actual")
plt.subplot(3, 1, 3)
plt.plot(X.iloc[:,3],pred,color='r',label="Predicted")
plt.scatter(X.iloc[:,3],Y,color='b',label="Actual")
plt.legend()
```

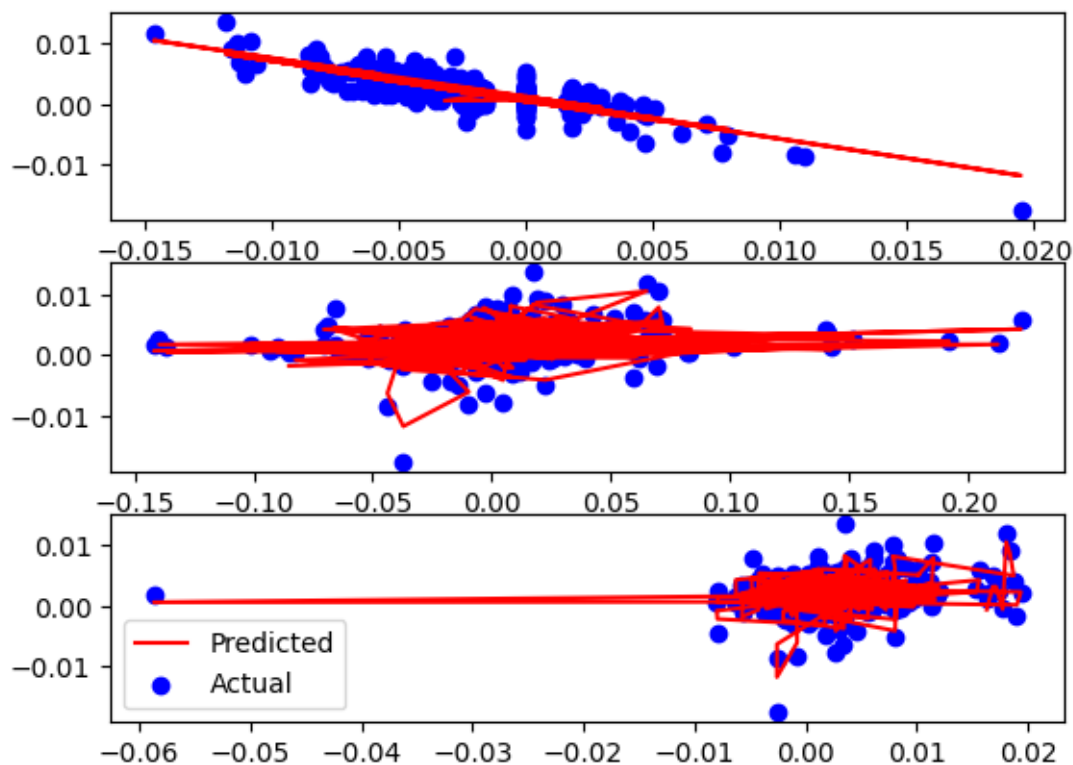
```
[168]: x=data1.loc[:
        ↪,['intercept','US_Dollar_Purchasing_power_per_change','CPI_Energy_per_change','Monthly_Hous
        x.iloc[:,3]
```

```
[168]: 0      0.000000
      1     -0.058703
      2     -0.004351
      3     -0.004370
      4      0.008413
      ...
     307     0.017422
     308     0.017127
     309     0.016838
     310     0.016556
     311     0.016290
      Name: Monthly_Housing_Cost_per_change, Length: 312, dtype: float64
```

```
[169]: Linear3(data1.loc[:
        ↪,['intercept','US_Dollar_Purchasing_power_per_change','CPI_Energy_per_change','Monthly_Hous
        ↪loc[:,['CPI_all_items_per_change']])
```

```
SSE: CPI_all_items_per_change    0.000894
dtype: float64
```

```
<ipython-input-167-f6a924f5ccfc>:12: MatplotlibDeprecationWarning: Auto-removal
of overlapping axes is deprecated since 3.6 and will be removed two minor
releases later; explicitly call ax.remove() as needed.
    plt.subplot(3, 1, 1)
```



4. Rent of Primary residence per change (CPI All items per change vs [US Dollar Purchasing Power per change, CPI Energy per change, Monthly Housing Cost per change, Rent of Primary residence per change]) In each step, we repeat the experiment with one extra column added. Plot the original data and the regression fit, and report the SSE in each linear regression experiment. Comment on which variables are most relevant in predicting the CPI All items per change based on the linear regression experiments performed.

```
[170]: def Linear4(X,Y):
    X_T=X.T
    B = np.linalg.inv(X_T @ X) @ X_T @ Y
    B.index=X.columns
    # print(B)
    pred=X @ B
    B.index=X.columns

    # print(X)
    print("SSE:",((Y - pred) ** 2).sum())
    sns.regplot(x=Y,y=pred,ci=None,color='red')
    plt.subplot(4, 1, 1)
    plt.plot(X.iloc[:,1],pred,color='r',label="Predicted")
    plt.scatter(X.iloc[:,1],Y,color='b',label="Actual")
    plt.subplot(4, 1, 2)
```

```

plt.plot(X.iloc[:,2],pred,color='r',label="Predicted")
plt.scatter(X.iloc[:,2],Y,color='b',label="Actual")
plt.subplot(4, 1, 3)
plt.plot(X.iloc[:,3],pred,color='r',label="Predicted")
plt.scatter(X.iloc[:,3],Y,color='b',label="Actual")
plt.subplot(4, 1, 4)
plt.plot(X.iloc[:,4],pred,color='r',label="Predicted")
plt.scatter(X.iloc[:,4],Y,color='b',label="Actual")
plt.legend()

```

```


[171]: Linear4(data1.loc[:
↳,['intercept','US_Dollar_Purchasing_power_per_change','CPI_Energy_per_change','Monthly_Hous
↳loc[:,['CPI_all_items_per_change']])

```

SSE: CPI_all_items_per_change 0.000894
dtype: float64

<ipython-input-170-e4ffc7faacc8>:12: MatplotlibDeprecationWarning: Auto-removal of overlapping axes is deprecated since 3.6 and will be removed two minor releases later; explicitly call ax.remove() as needed.

```
plt.subplot(4, 1, 1)
```



Because multiple linear regression makes hyper plane we are seeing a squiggle in the graph which

represents, for same value of x1 input variable we are getting different value of y as value of input variable x2 changes

when input variables are : 'US_Dollar_Purchasing_power_per_change','CPI_Energy_per_change','Monthly_Housing_Cost' the SSE is least which is 0.000894

```
[184]: import pandas as pd
import numpy as np

# Load the dataset into a Pandas DataFrame
df = pd.read_csv("/content/drive/My Drive/ProbStatsProject/CleanedDallas97-22.
↪csv")
df['DATE'] = pd.to_datetime(df['DATE']).dt.year

# Select the columns "Monthly Housing" and "All Items in the Urban Center" for
↪the years 1972-1996
data = df[(df['DATE'] >= 1997) & (df['DATE'] <= 2021)][["Monthly_Housing_Cost",
↪"CPI_all_items", "DATE"]]
```

```
[173]: # Fill in missing values with the average value for that year
data = data.groupby('DATE').transform(lambda x: x.fillna(x.mean()))
```

```
[174]: data.isna().sum()
```

```
[174]: Monthly_Housing_Cost    0
CPI_all_items                0
dtype: int64
```

```
[175]: # Categorize the data into high and low using the median value as a threshold
median_housing = np.median(data["Monthly_Housing_Cost"])
median_all_items = np.median(data["CPI_all_items"])
data["Category_Monthly_Housing_Cost"] = np.where(data["Monthly_Housing_Cost"] >
↪median_housing, "high", "low")
data["Category_CPI_all_items"] = np.where(data["CPI_all_items"] >
↪median_all_items, "high", "low")
```

```
[176]: # Create a contingency table
contingency_table = pd.crosstab(data["Category_Monthly_Housing_Cost"],
↪data["Category_CPI_all_items"])
```

```
[177]: data[data.Category_CPI_all_items=='low'].shape
```

```
[177]: (150, 4)
```

```
[178]: contingency_table
```

```
[178]: Category_CPI_all_items      high  low
Category_Monthly_Housing_Cost
```

high	135	15
low	15	135

```
[179]: # Perform chi-square analysis and calculate p-value manually
```

```
observed = contingency_table.values
row_totals = observed.sum(axis=1)
column_totals = observed.sum(axis=0)
total = observed.sum()
expected = np.outer(row_totals, column_totals) / total
chi2 = np.sum((observed - expected)**2 / expected)
df = (len(row_totals) - 1) * (len(column_totals) - 1)
```

```
[180]: chi2
```

```
[180]: 192.0
```

P-Value is calculated manually taking degree of freedom = 1 and chi square statistic 192 comes out to be less than 0.00001. $p\text{-value} < 0.05$ so we reject the null hypothesis so monthly housing and all item are not independent

3.6 Time Series Analysis

```
[180]:
```

```
[181]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

def cal_MSE(output):
    return round(sum((output[:,0]-output[:,1])**2)/len(output), 3)

def cal_mape(output):
    true = output[:,0]
    pred = output[:,1]
    return np.round(np.mean(np.abs((true - pred)/true))*100, 3)

df = pd.read_csv("/content/drive/My Drive/ProbStatsProject/CleanedDallas97-22.
↳ csv")
df = df.iloc[:, 1:]
df['DATE'] = pd.to_datetime(df.DATE)

data = df[(df['DATE'] >= '2018-01-01') & (df['DATE'] <= '2021-12-01')]

train_frame = data[(data['DATE'] >= '2018-01-01') & (data['DATE'] <=
↳ '2020-12-01')][['DATE', 'Rent_of_Primary_residence']]
test_frame = data[(data['DATE'] >= '2021-01-01') & (data['DATE'] <=
↳ '2021-12-01')][['DATE', 'Rent_of_Primary_residence']]
```



```
train = train_frame['Rent_of_Primary_residence'].tolist()
test = test_frame['Rent_of_Primary_residence'].tolist()
```

EWMA

[181]:



```
[182]: def ewma(alpha, col):
        ewma = [col[0]]
        for i in range(1, len(col)):
            ewma.append(alpha * col[i] + (1 - alpha) * ewma[i-1])
        return ewma

def pred_ewma(model, alpha, col):
    prediction = [train[-1]]
    for i in range(0, len(col)):
        prediction.append(alpha * col[i] + (1 - alpha) * prediction[-1])
    return np.array(prediction[1:])

def run_ewma(alpha, train, test):
    model = ewma(alpha, train)
    pred = pred_ewma(model, alpha, test)

    output = np.array([[i[0], i[1]] for i in zip(test, pred)])
    plt.figure(figsize=(14, 5))
    plt.plot(data['DATE'], train+output[:, 0].tolist(), linewidth=2,
    ↪color='Black')
    plt.plot(test_frame['DATE'], output[:, 1], linewidth=2, color='red')
    plt.title("Alpha: "+str(alpha)+ ", MSE = "+str(cal_MSE(output))+ " & MAPE =_
    ↪"+str(cal_mape(output)))
    plt.show()

alpha1=0.5
alpha2=0.8
run_ewma(alpha1, train, test)
run_ewma(alpha2, train, test)
```



4 Auto-Regression

[182]:

[183]: `def get_feature(p, data, pred=False, init=None):`

```

X = []
Y = []
if pred:
    X.append(init)
    Y.append(data[0])
    for i in range(1, p):
        tmp = X[-1]
        tmp = tmp[1:]
        tmp.append(Y[-1])
        X.append(tmp)

```

```

        Y.append(data[i])

    for i in range(p, len(data)):
        X.append([data[i - idx] for idx in range(1, p+1)][::-1])
        Y.append(data[i])
    X = np.array(X)
    Y = np.array(Y)
    return X, Y

def AutoRegression(p, data):
    X, Y = get_feature(p, data, False)
    beta_cap = np.dot(np.dot(np.linalg.inv(np.dot(X.T, X)), X.T), Y)
    return beta_cap

def AutoRegPred(p, beta_cap, data, init):
    X, Y = get_feature(p, data, True, init)
    pred = np.dot(beta_cap, X.T)
    return Y, pred.tolist()

def run_auto_regression(p, train, test):
    true, pred = AutoRegPred(p, AutoRegression(p, train), test, train[-p:])
    output = np.array([[i[0], i[1]] for i in zip(true, pred)])

    plt.figure(figsize=(14, 5))
    plt.plot(data['DATE'], train+output[:, 0].tolist(), linewidth=2,
    color='Black')
    plt.plot(test_frame['DATE'], output[:, 1], linewidth=2, color='red')
    plt.title("P: "+str(p)+ ", MSE = "+str(cal_MSE(output))+" & MAPE = "+str(cal_mape(output)))
    plt.show()

p1 = 3
p2 = 10
run_auto_regression(p1, train, test)
run_auto_regression(p2, train, test)

```

